

eda

June 1, 2020

```
[1]: import pandas as pd
import numpy as np
import plotly.express as px
import plotly.graph_objects as go
clrs = ['#025159', '#03A696', '#F28705', '#F25D27', '#F20505']
```

0.1 # Explorative Data Analysis for Challenge FS20C7 - Recommendations im Detailhandel

Author Roman Studer

This notebook is dedicated to the explorative data analysis of the data set used in the [Challenge FS20C7 - Recommendations in Retail](#). Since the execution of the calculations to create the graphics in this notebook is rather long, a PDF and HTML version is available here. Please note that in the PDF the interactivity of the graphics is lost.

0.2 1. Dataset

Import and first look at Dataset

```
[ ]: df = pd.read_csv('Recommender4Retail.csv', index_col='Unnamed: 0');
```

```
[3]: df.head()
```

```
[3]:
```

	order_id	user_id	eval_set	order_number	order_dow	order_hour_of_day	\
1	2539329	1	prior	1	2	8	
2	2539329	1	prior	1	2	8	
3	2539329	1	prior	1	2	8	
4	2539329	1	prior	1	2	8	
5	2539329	1	prior	1	2	8	

	days_since_prior_order	product_id	add_to_cart_order	reordered	\
1	NaN	196	1	0	
2	NaN	14084	2	0	
3	NaN	12427	3	0	
4	NaN	26088	4	0	
5	NaN	26405	5	0	

	product_name	aisle_id	department_id	\
1	Soda	77	7	
2	Organic Unsweetened Vanilla Almond Milk	91	16	
3	Original Beef Jerky	23	19	
4	Aged White Cheddar Popcorn	23	19	
5	XL Pick-A-Size Paper Towel Rolls	54	17	

	department	aisle
1	beverages	soft drinks
2	dairy eggs	soy lactosefree
3	snacks	popcorn jerky
4	snacks	popcorn jerky
5	household	paper goods

```
[4]: df.tail()
```

```
[4]:
```

	order_id	user_id	eval_set	order_number	order_dow	\
33819102	272231	206209	train	14	6	
33819103	272231	206209	train	14	6	
33819104	272231	206209	train	14	6	
33819105	272231	206209	train	14	6	
33819106	272231	206209	train	14	6	

	order_hour_of_day	days_since_prior_order	product_id	\
33819102	14	30.0	40603	
33819103	14	30.0	15655	
33819104	14	30.0	42606	
33819105	14	30.0	37966	
33819106	14	30.0	39216	

	add_to_cart_order	reordered	\
33819102	4	0	
33819103	5	0	
33819104	6	0	
33819105	7	0	
33819106	8	1	

	product_name	aisle_id	department_id	\
33819102	Fabric Softener Sheets	75	17	
33819103	Dark Chocolate Mint Snacking Chocolate	45	19	
33819104	Phish Food Frozen Yogurt	37	1	
33819105	French Baguette Bread	112	3	
33819106	Original Multigrain Spoonfuls Cereal	121	14	

	department	aisle
33819102	household	laundry
33819103	snacks	candy chocolate

```

33819104    frozen    ice cream ice
33819105     bakery          bread
33819106  breakfast          cereal

```

```
[5]: df.columns
```

```
[5]: Index(['order_id', 'user_id', 'eval_set', 'order_number', 'order_dow',
          'order_hour_of_day', 'days_since_prior_order', 'product_id',
          'add_to_cart_order', 'reordered', 'product_name', 'aisle_id',
          'department_id', 'department', 'aisle'],
          dtype='object')
```

0.3 2. Data description

The following columns are included in the dataset: - `order_id`: order identifier - `user_id`: customer identifier - `eval_set`: which evaluation set this order belongs in (see SET described below) - `order_number`: the order sequence number for this user (1 = first, n = nth) - `order_dow`: the day of the week the order was placed on - `order_hour_of_day`: the hour of the day the order was placed on - `days_since_prior`: days since the last order, capped at 30 (with NAs for `order_number` = 1) - `product_id`: product identifier - `product_name`: name of the product - `aisle_id`: aisle identifier - `aisle`: the name of the aisle - `department_id`: department identifier - `department`: the name of the department - `reordered`: 1 if this product has been ordered by this user in the past, 0 otherwise

Where SET is one of the four following evaluation sets (`eval_set` in orders):

- “prior”: orders prior to that users most recent order
- “train”: training data supplied to participants
- “test”: test data reserved for machine learning competitions

```
[6]: # shape
      f'Dimensionality of the dataset {df.shape}'
```

```
[6]: 'Dimensionality of the dataset (33819106, 15)'
```

0.4 3. Data cleaning

Dropping columns and deal with missing data

0.4.1 3.1. Dropping columns

The following columns are not included in the analysis: - `eval_set`: for the further use of the dataset a categorization in “prior”, “train” or “test” of the data is not necessary. - `order_number`: the order of orders per user is not included. This information can also be read from the ‘`order_id`’. - `add_to_cart_order`: the order of how a product has been placed is not relevant for the planned recommender - `order_dow` and `order_hour_of_day`: time of order is irrelevant for the planned recommender - `days_since_prior_order`, also does not flow into the planned recommender.

```
[7]:
```

```
df = df.  
↳ drop(['eval_set', 'order_number', 'add_to_cart_order', 'order_dow', 'order_hour_of_day',  
↳ 'days_since_prior_order'], axis=1)
```

0.4.2 3.2. Data types

```
[8]: # checking datatypes  
df.dtypes
```

```
[8]: order_id          int64  
     user_id          int64  
     product_id       int64  
     reordered        int64  
     product_name     object  
     aisle_id         int64  
     department_id    int64  
     department       object  
     aisle            object  
     dtype: object
```

Data types per feature are correct.

0.4.3 3.3. Missing values

Rows with missing values can no longer be used. Since these are individual transactions that at first glance do not appear to have any connection (apart from possible customer preferences) we do not impute missing data.

```
[9]: df.drop_duplicates(inplace=True)  
df.isna().sum()
```

```
[9]: order_id          0  
     user_id          0  
     product_id       0  
     reordered        0  
     product_name     0  
     aisle_id         0  
     department_id    0  
     department       0  
     aisle            0  
     dtype: int64
```

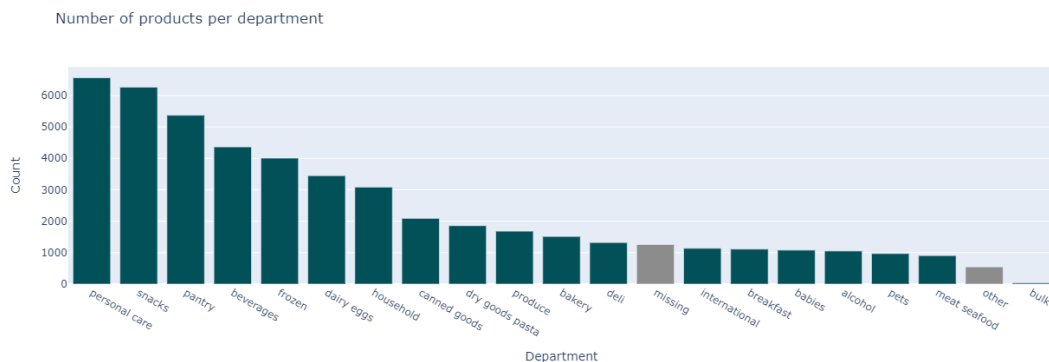
0.5 4. Products

Visual analysis of the distribution of products in the dataset.

```
[10]: df_department = df.groupby(by='department')['product_id'].nunique().
      ↪reset_index()

color = [clrs[0],]*len(df_department.department.unique())
color[14], color[15] = '#8c8c8c', '#8c8c8c'
fig = go.Figure(data=[go.Bar(
    x=df_department.department,
    y=df_department.product_id,
    marker_color=color, # marker color can be a single color value or an
    ↪iterable
)])
fig.update_xaxes(categoryorder='total descending')
fig.update_layout(title_text='Number of products per department',
                  xaxis_title="Department",
                  yaxis_title="Count")

fig.show()
```



The products are divided into 19 departments. In the graphic above, the departments Pantry, Personal Care, Snacks, Beverages and Frozen are particularly prominent. It can also be seen that there are products that cannot be assigned to a specific department. These 548 products fall into the category ‘other’. A larger proportion (1258 products) cannot even be assigned to the category ‘other’ and were given the designation ‘missing’. ‘missing’ and ‘other’ are greyed out in the graphic.


```
[11]: # products per aisle
df_aisle = df.groupby('aisle')['product_id'].nunique().reset_index()

fig = px.bar(df_aisle,
             x='aisle',
             y='product_id',
             title='Number of products per aisle',
             labels = {'product_id':'Count', 'aisle':'Aisle'},
```

```

color_discrete_sequence = [clrs[0]]).
↪update_xaxes(categoryorder='total descending')
fig.show()

```



The individual aisles of the departments, listed by number of products, show a range from 1258 up to 12 products per aisle.

```

[12]: df_dep_aisle = df.groupby(by=['department', 'aisle'])['product_id'].nunique().
↪reset_index()
fig = px.treemap(df_dep_aisle,
                 path=['department', 'aisle'],
                 values='product_id',
                 title='Ratio of orders in departments and aisles',
                 color='department')
fig.show()

```

Ratio of orders in departments and aisles



Product groups (Aisles) in the departments can be visualized well in a treemap. It is well visible that the top 5 departments contain slightly more than half of the available product range. To focus on a department, click on it.

0.5.1 Most ordered products

```
[13]: # most ordered products
most_ordered = df['product_name'].value_counts()
# top 20 products
print(f'Top 20 Products: \n{most_ordered[:20]}')
```

```
Top 20 Products:
Banana                491291
Bag of Organic Bananas 394930
Organic Strawberries  275577
Organic Baby Spinach  251705
Organic Hass Avocado  220877
Organic Avocado       184224
Large Lemon          160792
Strawberries          149445
Limes                 146660
Organic Whole Milk    142813
Organic Raspberries   142603
Organic Yellow Onion   117716
Organic Garlic         113936
Organic Zucchini       109412
Organic Blueberries    105026
Cucumber Kirby        99728
Organic Fuji Apple     92889
Organic Lemon          91251
Organic Grape Tomatoes 88078
Apple Honeycrisp Organic 87272
Name: product_name, dtype: int64
```

At first glance, the top products appear to be organic products such as fruits and vegetables. At the very top, the banana.


0.6 Orders

Visual analysis of the distribution of orders in the dataset.

```
[14]: # Number of products in orders
df_n_order = df.groupby(by=['order_id'])['product_id'].count().reset_index()
n_orders = df_n_order.product_id.value_counts()

fig = px.bar(x=n_orders.index,
              y=n_orders.values,
              title='Number of products in a shopping cart',
              labels={'x': 'Number of products in basket', 'y': 'Number of orders'},
              color_discrete_sequence=[clrs[1]])

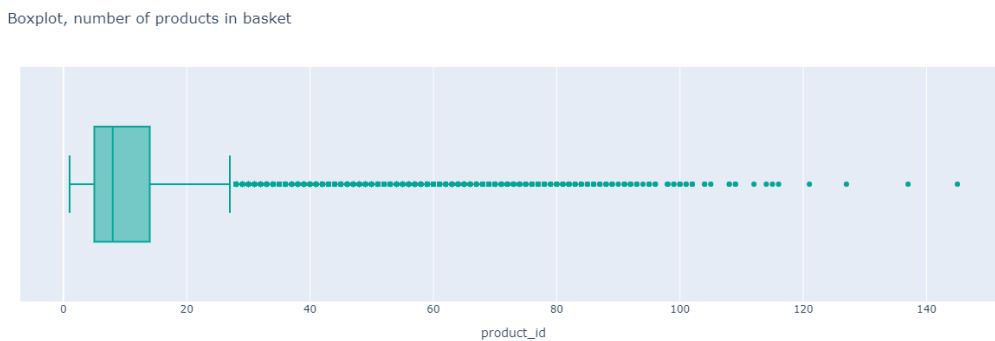
fig.show()
```



The upper graphic shows how often a shopping cart contains a certain number of products. You can see here that there are orders with more than 140 products. The vast majority, however, have between one and twenty products in their shopping cart. The peak lies at five products with over 237 thousand orders containing this number of products.

```
[15]: # orders per customer
df_c_order = df.groupby(by=['order_id'])['product_id'].count().reset_index()
fig = px.box(data_frame=df_c_order,
             x='product_id',
             title='Boxplot, number of products in basket',
             labels={'x': 'Number of products in order'},
             color_discrete_sequence=[clrs[1]],
             orientation='h')

fig.show()
```




The box plot confirms the image of the upper bar plot. The minimum size of an order is one, since a smaller number would not trigger a transaction. The maximum size of an order is 145 and is an exception. The median is eight products, with 50% of orders containing between 5 and 14 products.


```
[16]: # number of orders per customer
df_n_customer = df.groupby(by='user_id')['order_id'].count().reset_index()

fig = px.histogram(data_frame=df_n_customer,
                   x="order_id",
                   title='Number of customers with x orders',
                   labels={'order_id': 'Number of orders', 'count': 'Amount of_
→users'},
                   color_discrete_sequence=[clrs[1]],
                   nbins=150
                   )

fig.update_layout(yaxis_title="Amount of users")

fig.show()
```



The distribution of the number of orders has shifted strongly to the left. The histogram therefore moves to the lower left limit. The graphic shows how many customers (y-axis) carry out a certain number of orders (x-axis). If we look at the bar on the far left, we can see that almost 60 thousand customers placed between 1 and 50 orders. On the far right of the plot, barely visible, the highrollers with over 3500 orders.

```
[17]: # number orders per customer
fig = px.box(data_frame=df_n_customer,
             x='order_id',
             title='Boxplot, Number of orders from customers',
             labels={'order_id': 'Number of orders'},
             color_discrete_sequence=[clrs[1]],
             orientation='h')

fig.show()
```

Boxplot, Number of orders from customers



The box plot to the upper histogram shows that 50% of the users place between 44 and 196 orders. The median is at 90.

1 Conclusion

The insights gained from this brief analysis are incorporated into the recommender system. Above all, the consequences of various types of product reduction are well illustrated here. For example, if I remove aisles with few products, I reduce the number of products only slightly, but lose a large part of the product variety. It is also interesting to see that the median number of orders per user is 90.